

Introduction

INF8422: Robotic Perception and Spatial Intelligence

Prof. Pierre-Yves Lajoie

**POLYTECHNIQUE
MONTREAL**


Agenda

1. Course Logistics
2. What is robot perception and spatial intelligence?
3. Semester Overview
4. Autonomous Robotics
5. 3D Geometry
6. LiDAR Sensor
7. Iterative Closest Point (ICP)


Prerequisites



Linear algebra and basic calculus



Python programming



Patience and resourcefulness

This is the first iteration of the course.

Help us make it converge: share your feedback (and your gradients).

Pierre-Yves Lajoie eng. Ph.D.

Assistant Professor - Department of Computer and Software Engineering

MIST Laboratory: mistlab.ca



Autonomous robotics research laboratory

Spatial Intelligence, Localization, Mapping, Planning, Multi-agent Systems, etc.

Course Logistics

Assessment

Assessment	Weight
5 Lab Assignments	$5 \times 5\% = \mathbf{25\%}$
4 best Quizzes / 5	$4 \times 7.5\% = \mathbf{30\%}$
Course Project	25%
Final Exam	20%

In-class quiz

- About every two weeks, **25 min**, start of class.
- Covers material from previous lectures.
- 5 quizzes total, **only the 4 best** count.
- First quiz: **September 11**.

Double threshold: You must have a minimum average of **50% on individual assessments** (Quizzes and Exam) to obtain a passing grade.

Labs

Organization

- **5 labs**, individually or in teams of 2.
- Lab instructor: **Julien Lagacé**.
- All resources on **Moodle**.
- Join the course **Discord** for support (link on Moodle).

Labs

Lab sessions start on September 9.

First submission:

Lab 1: Wednesday, **September 15**, 11:30 PM.

Session project

Format

- Team of **2 to 3 people**.
- Report of **4 to 6 pages**, double column, **IEEE** format.
- A guide for the project is available on **Moodle**.
- Criteria: scientific rigor, experiments, interpretation of results.

Schedule

October 9, 8:30 AM – Team composition + topic description (max. 1 page): objective (What?), motivation (Why?), and planned experiments (How?).


November 27 – **10-minute** presentation.

December 22, 8:30 AM – Final report submission.

Artificial intelligence - Students' point of view

Impact on your learning

- Hinders skill development.
- Decline in critical thinking.
- Limited creativity.



Your constraints

- You want good grades.
- You don't have much time (courses, research, etc.).
- Sometimes what is asked is unrealistic without AI.

Artificial intelligence - Teacher's point of view

Major problem: How to design labs that allow understanding the concepts seen in class **AND** are complex enough to outsmart AI...

Solution 1 - Make the labs harder

Positive:

- Allows more ambitious labs.
- Closer to the work context.
- Learn to use AI tools.

Negative:

- You have no choice but to use them, and it limits learning.

Solution 2 - Make the labs easier

Proposed in *Manifesto for Human Education*, J. Pilon, 2026

Positive:

- Can be done without AI
- Maximizes learning

Negative:

- Can be solved by AI and hard to evaluate

Artificial Intelligence - Instructor's Point of View

Solution 1 - Term Project

Term project:

You can use AI if you want and make the project as ambitious as possible.

That said, if you wish *not to use AI* for the project, please let us know, we will adjust our expectations accordingly.

Solution 2 - Labs

We ask you not to use AI in the labs, for the sake of your learning.

In return:

- Support from the teaching team
- Adjustments if the lab is too difficult or too long

Before using AI to solve the lab, come see us and ask for help.

We can explain, help you, and adjust if necessary.

Session Calendar

September

Sep 11 – Quiz 1

Sep 15 – Lab 1 due

Sep 23-25 – Remote lab and course

Sep 29 – Lab 2 due

October

Oct 2 – Quiz 2

Oct 9 –  Team + topic

Oct 20 – Lab 3 due

Oct 23 – Quiz 3

November

Nov 3 – Lab 4 due

Nov 6 – Quiz 4

Nov 17 – Lab 5 due

Nov 20 – Quiz 5

Nov 27 –  Presentation

December

Dec 22 –  Report due

Legend

 Quiz

 Lab submissions

 Project

What is robot perception and spatial intelligence?

Myth or Reality?

"robotic perception == computer vision."

Perception $\neq$ Vision

Partially false. Computer vision deals with images and videos: object detection, segmentation, shape recognition, etc.

Robots need more information:

- **3D geometry of the entire environment:** depth, obstacles, traversability, structures.
- **Physical properties:** articulation, mass, forces.
- **Real-time constraints:** an error = a collision.

Vision: "What is it?"

→ Classification, segmentation, etc.

Robotic perception:

→ Where am I?

→ How do I get to point B without colliding?

→ How do I pick up an object?

→ What will the state of the environment be after an action?

Myth or Reality?

"Robotic perception is harder than computer vision."

It's all relative

Sometimes yes:

- Must produce richer outputs (3D, physical properties, etc).
- Sensitive to processing speed.
- Errors have real consequences.

Sometimes no: the robot can act to simplify perception.

The robot is not passive

Unlike classical vision, the robot can:

- Move to find a better viewing angle.
- Get closer to improve resolution.
- Interact to reveal hidden properties.

Myth or Reality?

"Robotic perception is limited to visual data."

Robotic perception is multimodal

False. Each modality provides complementary information.

Vision

RGB, RGB-D, stereo camera

Active geometry

LiDAR, Radar, Sonar

Inertial

IMU (accelerometer, gyroscope, magnetometer)

Tactile

Force, pressure, texture

Infrastructure

GPS, UWB, WiFi

Acoustic

Microphones

Myth or Reality?

"Robotic perception == state estimation."

- Table 1**



Semester Overview

The course covers several topics:

1. **Foundations:** 3D geometry and mathematical tools
2. **SLAM (Simultaneous Localization and Mapping):** Sensor data extraction, state estimation.
3. **Representation:** How to model the world and handle uncertainty.
4. **Dynamic Perception and Action:** Dynamics, interaction, manipulation, multi-agent.
5. **Applications:** Advanced topics and case studies.



Autonomous Robotics


The Autonomy Cycle: Sense-Plan-Act

The classic paradigm of autonomous robotics (Sense-Plan-Act):

- **Sense:** Data acquisition and estimation of the state of the world.
- **Predict:** Predict the evolution of the environment.
- **Plan:** Decision-making based on the estimated state to achieve a goal.
- **Act:** Execution of motor commands. Changes the state of the robot and the environment.

Role of perception

This is the interface between the real physical world and the robot's digital model of itself and its environment.



Paradigm 1: Perception for Autonomy

Historically, perception has been seen as serving navigation.

- **Objective:** Enable the robot to move without collision and to localize itself.
- **Classical approach:** The robot creates and maintains a map of its environment in real time. The challenge is to find the map and position estimate that best explain the noisy sensor data.

Typical Applications

- Autonomous vehicles (Waymo, Tesla).
- Robot vacuums (iRobot).
- Logistics (Warehouse management, Amazon).

Paradigm 2: Autonomy for Perception

Active Perception: The robot acts to perceive better.

- **Example:** "Moving to see behind the obstacle" or "Approaching to classify better".
- **Objective:** Maximizing information, map quality, etc.

Typical Applications

- Infrastructure inspection.
- Search and rescue.
- Exploration of unknown environments.
- Precision agriculture.

3D Geometry

3D Geometry

Robotic navigation relies on 3D geometry:



- **Coordinate Frames** (Frames)
- **Positions and Translations**
- **Orientation Representation**
- **Pose Representation** (Position + Orientation)

Coordinate Frames (Coordinate Frames)

A frame is a set of orthogonal axes attached to a body.

Common frames in robotics:

- **Robot frame (r):** Attached to the robot (often at the center of mass).
- **Sensor frame ($c, l \dots$):** Attached to the sensor (e.g., camera, LiDAR).
- **World frame (w):** Fixed in the environment (e.g., starting point).





Right-Handed Convention

$$\mathbf{x} \times \mathbf{y} = \mathbf{z}$$

Mnemonics

- Thumb ($\mathbf{z}$), Index ($\mathbf{x}$), Middle finger ($\mathbf{y}$).
- Thumb toward $\mathbf{z}$, fingers curl from $\mathbf{x}$ to $\mathbf{y}$.



Common frames (left: robot frame, right: camera/image frames)

Frame Conventions

Robot Frame (r)



- Origin: Center of mass.
- x_r : Forward.
- y_r : Left.
- z_r : Up.

Image Frame (i) (2D)

- Origin: Top-left corner.
- x right, y down.

Camera Frame (c)

- Origin: Optical center.
- z_c : Looks at the scene (Optical Axis).
- x_c : To the right.
- y_c : Downward.




Points and Positions

The position of a point p in the frame w is a vector:

$$p^w = \begin{bmatrix} p_x^w \\ p_y^w \\ p_z^w \end{bmatrix} \in \mathbb{R}^3$$

The scalars p_x^w, p_y^w, p_z^w are the projections onto the axes x_w, y_w, z_w .



Translations and Composition

Translation: between two points p_1 and p_2 :

$$p_{12}^w = p_2^w - p_1^w$$


Composition:

$$p_2^w = p_1^w + p_{12}^w$$

Inverse:

$$p_{12}^w = -p_{21}^w$$

Note: The notation makes the frame (w , i.e. world) explicit as a superscript.



Rotation Representations

How do we describe the orientation of frame r (robot) relative to w (world)?

1. **Euler Angles** (Roll-Pitch-Yaw)
2. **Axis-Angle**
3. **Quaternions**
4. **Rotation Matrix**

1. Euler Angles (Roll-Pitch-Yaw)

An RPY rotation requires specifying the order of the axes.


Extrinsic XYZ convention (fixed axes):

$$R = R_z(\gamma) R_y(\beta) R_x(\alpha)$$

α : roll around x ; β : pitch around y ; γ : yaw around z .

Advantages

Intuitive. Minimal (3 parameters).




Euler Angles – Limitations

Gimbal Lock (Singularity)

If the pitch $\beta = \pm\pi/2$, one degree of freedom is lost: two rotation axes become collinear.

Other drawbacks:

- Costly trigonometric computations for composition and inversion.
- Several possible conventions (RPY, YPR, XYZ, ZYX...) → frequent sources of confusion.



2. Axis-Angle Representation

Euler's theorem: any rotation is equivalent to a rotation of angle θ about a fixed axis $\mathbf{u}$ ($\|\mathbf{u}\| = 1$).

The **rotation vector** $\boldsymbol{\omega} = \theta\mathbf{u} \in \mathbb{R}^3$ contains the three independent parameters.

Rodrigues' Formula (Axis-Angle $\rightarrow$ Matrix):

$$R_r^w = \cos(\theta) \mathbf{I}_3 + \sin(\theta)[\mathbf{u}]_{\times} + (1 - \cos(\theta)) \mathbf{u}\mathbf{u}^T$$

Where $[\mathbf{u}]_{\times}$ is the skew-symmetric matrix such that $[\mathbf{u}]_{\times} \mathbf{v} = \mathbf{u} \times \mathbf{v}$:

$$[\mathbf{u}]_{\times} = \begin{bmatrix} 0 & -u_z & u_y \\ u_z & 0 & -u_x \\ -u_y & u_x & 0 \end{bmatrix}$$

Matrix $\rightarrow$ Axis-Angle Conversion

How to recover $(\mathbf{u}, \theta)$ from a matrix R ?

Angle θ : via the trace of the matrix:

$$\text{tr}(R) = 1 + 2 \cos(\theta) \implies \theta = \arccos\left(\frac{\text{tr}(R) - 1}{2}\right)$$

Axis $\mathbf{u}$: eigenvector associated with the eigenvalue $\lambda = 1$:

$$R \mathbf{u} = \mathbf{u}$$

The axis of rotation is invariant under the rotation.

Eigenvalues of R

The eigenvalues of any rotation matrix are:

$$\{1, e^{i\theta}, e^{-i\theta}\}$$

3. Quaternions

Extension of complex numbers: $q = q_4 + i q_1 + j q_2 + k q_3$ with $i^2 = j^2 = k^2 = ijk = -1$, $ij = -ji = k$, $jk = -kj = i$, $ki = -ik = j$.

Convention: $q = \begin{bmatrix} v \\ w \end{bmatrix} = \begin{bmatrix} q_1 \\ q_2 \\ q_3 \\ q_4 \end{bmatrix}$,

Link with Axis-Angle ($\mathbf{u}, \theta$):

$$q = \begin{bmatrix} \mathbf{u} \sin(\theta/2) \\ \cos(\theta/2) \end{bmatrix}$$

with v the vector part and w the scalar part.

Link with the rotation matrix $R(q)$:

$$R(q) = \begin{bmatrix} q_1^2 - q_2^2 - q_3^2 + q_4^2 & 2(q_1 q_2 - q_3 q_4) & 2(q_1 q_3 + q_2 q_4) \\ 2(q_1 q_2 + q_3 q_4) & -q_1^2 + q_2^2 - q_3^2 + q_4^2 & 2(q_2 q_3 - q_1 q_4) \\ 2(q_1 q_3 - q_2 q_4) & 2(q_2 q_3 + q_1 q_4) & -q_1^2 - q_2^2 + q_3^2 + q_4^2 \end{bmatrix}$$

Unit Quaternion Operations

To represent a rotation: unit quaternions ($\|q\| = 1$), with $q = \begin{bmatrix} v \\ w \end{bmatrix}$.

Composition (Product):

$$q_a \otimes q_b = \begin{bmatrix} w_a v_b + w_b v_a + v_a \times v_b \\ w_a w_b - v_a^T v_b \end{bmatrix}$$

Expressing a point in a different frame:

$$q_r^w \otimes \begin{bmatrix} p^r \\ 0 \end{bmatrix} \otimes (q_r^w)^{-1} = \begin{bmatrix} p^w \\ 0 \end{bmatrix} = \begin{bmatrix} R(q_r^w) p^r \\ 0 \end{bmatrix}$$

Inverse:

$$q^{-1} = \begin{bmatrix} -q_{1:3} \\ q_4 \end{bmatrix}$$

Advantages of Quaternions

Advantages

- **No singularity** – no Gimbal Lock.
- **Compact** – 4 scalars and a normalization constraint, so 3 degrees of freedom.
- **Efficient computation** – composition requires no trigonometry: 16 multiplications vs 27 for matrices.
- **Generation** – simply generate 4 real numbers and normalize to obtain a valid quaternion.

Drawback: Double Cover

q and $-q$ represent the **same rotation**.

This causes problems for interpolation and optimization (discontinuities).


4. Rotation Matrix

We project the axes of r (x_r, y_r, z_r) into the frame w and stack these vectors as columns.

$$R_r^w = \begin{bmatrix} | & | & | \\ x_r^w & y_r^w & z_r^w \\ | & | & | \end{bmatrix} \in SO(3)$$

2D example:

$$R_r^w = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \in SO(2)$$



Properties of the $SO(3)$ Group

Orthogonality

Unit and perpendicular columns:

$$(R_r^w)^T R_r^w = I_3 \implies (R_r^w)^{-1} = (R_r^w)^T$$

Right-Handed Frame

$$\det(R_r^w) = +1$$

Scalar triple product of the columns: $(x \times y) \cdot z = 1$.

Operations with Rotation Matrices

Point rotation: convert p^r (robot frame) to p^w (world frame):

$$p^w = R_r^w p^r$$

Composition: rotation of the sensor in the world frame:


$$R_c^w = R_r^w R_c^r$$

Non-commutative

$$R_r^w R_c^r \neq R_c^r R_r^w$$

Inverse:

$$R_r^w = (R_w^r)^{-1} = (R_w^r)^T$$



Poses and Rigid Transformations

Pose and Rigid Transformation


A **Pose** combines the position and orientation of r relative to w :

$$(R_r^w, t_r^w)$$

- R_r^w : Orientation of robot r in the world frame w .
- t_r^w : Position of the robot's frame r in the world frame w .

Point transformation (change of frame):

$$p^w = R_r^w p^r + t_r^w$$



Homogeneous Coordinates

Homogeneous coordinates allow an affine transformation to be represented by a 4×4 matrix multiplication.

Homogeneous point: $\tilde{p}^r = \begin{bmatrix} p^r \\ 1 \end{bmatrix}$

Transformation Matrix $T_r^w \in SE(3)$:

$$T_r^w = \begin{bmatrix} R_r^w & t_r^w \\ 0_{1 \times 3} & 1 \end{bmatrix}$$

The frame change equation becomes:

$$p^w = R_r^w p^r + t_r^w \quad \Longleftrightarrow \quad \tilde{p}^w = T_r^w \tilde{p}^r$$

$SE(3)$ Pose Operations

Transformation of a point:

$$\begin{bmatrix} p^w \\ 1 \end{bmatrix} = \begin{bmatrix} R_r^w & t_r^w \\ 0 & 1 \end{bmatrix} \begin{bmatrix} p^r \\ 1 \end{bmatrix}$$


Composition (common intermediate frame):

$$T_c^w = T_r^w T_c^r$$

Inverse:

$$T_w^r = (T_r^w)^{-1} = \begin{bmatrix} (R_r^w)^T & -(R_r^w)^T t_r^w \\ 0 & 1 \end{bmatrix}$$

Note: $(T_r^w)^{-1} \neq (T_r^w)^T$ because the matrix T is not orthogonal.



Summary of Representations

Representation	Scalars (constraints)	Singularities	Main Use
Rotation Matrix	9 (6)	No	Operations on vectors
Euler Angles (RPY)	3	Yes	Intuitive display
Axis-Angle	3 (ω)	Yes, at $\theta = 0$	Calculations, conversions
Quaternion	4 (1)	No	Storage, generation

Generally for

- **Store / generate** → Quaternion
- **Compute / apply** → Rotation Matrix or Axis-Angle
- **Communicate to a human** → Euler Angles

LiDAR Sensor

Sensor Taxonomy

Proprioceptive (Internal)

Measures the internal state of the system.

- Ex: Wheel encoders, IMU, Battery.

Passive

Receives energy. Consumes less energy.

- Ex: Cameras

Exteroceptive (External)


Measures the environment.

- Ex: Cameras, LiDAR, Radar.

Active

Emits energy. Consumes a lot of energy. Generally more robust.

- Ex: LiDAR, Sonar, Radar



Time-of-Flight (ToF) Sensors

Basic principle for LiDAR, Radar, Sonar.

$$d = \frac{c \cdot \Delta t}{2}$$

- c : Wave speed ($3 \cdot 10^8$ m/s for light, **343** m/s for sound).
- Δt : Round-trip time.

Comparison

- **LiDAR**: Collimated beam (Laser). Very precise pointwise measurement.
- **Sonar**: Emission cone. Measures the first reflection in the cone (angular ambiguity).


3D LiDAR and Point Clouds

Combination of ToF distance + scanning mechanism (rotating mirror).

- **Output:** Point cloud $\mathcal{P} = \{p_1, \dots, p_N\}$ with $p_i = (x, y, z, \text{reflectivity})$.
- **Spherical geometry:**

$$p_i = \begin{bmatrix} x_i \\ y_i \\ z_i \end{bmatrix} = r \begin{bmatrix} \cos \phi \sin \theta \\ \cos \phi \cos \theta \\ \sin \phi \end{bmatrix}$$

where r is the distance, θ the azimuth (rotation), ϕ the elevation (laser ID).




A LiDAR Problem: Motion Distortion

A scan (rotation of the laser) is not instantaneous (e.g., 10 Hz $\rightarrow$ 100 ms).

The Phenomenon

Point p_0 is acquired at t . Point p_N is acquired at $t + 100$ ms.

- In the meantime, the robot may have moved forward 2 m (at a speed of 72 km/h).
- The robot may have made a fast rotation or experienced vibrations.
- Result: Straight walls become curved, objects are duplicated.



Left: distorted scan. Middle: correction. Right: corrected scan (Deskew).

Deskewing – Motion Distortion Correction


The Solution: Deskewing

We project all points to a reference time t_{ref} (start or end of scan).

$$p_{\text{corrected}} = T_{\text{ref}}^{-1} \cdot T(t_i) \cdot p_{\text{raw}}$$

Requires knowing the robot's velocity at high frequency (IMU) to interpolate $T(t_i)$.

A scan is not a snapshot




Iterative Closest Point (ICP)

Point Cloud Alignment

Fundamental algorithm for autonomous robotics.

- Estimate the displacement of an autonomous car.
- Associate a point cloud with a CAD model.
- Guide a surgical tool
- etc.




Point Cloud Alignment

The Problem: Let two point clouds be given:

- **Target (Q):** The fixed cloud (e.g., the map).
- **Source (P):** The cloud to align (e.g., current scan).

Objective: Find the rigid transformation (R, t) that aligns P onto Q by minimizing the mean squared error.

$$E(R, t) = \sum_{i=1}^N \|q_i - (Rp_i + t)\|^2$$



The Chicken or the Egg?

The Dilemma

- If we knew the exact **correspondences**, we could compute (R, t) directly in a single step (closed-form solution).
- If we knew the **transformation** (R, t) , we could find the correspondences easily.

ICP solution: Proceed iteratively.

1. Estimate the correspondences (nearest neighbor).
2. Compute the optimal transformation for these correspondences.
3. Apply the transformation and start over.

The Standard ICP Algorithm

Input: P (Source), Q (Target), initial estimate (R_0, t_0) .


1. **Data association:** For each point $p_i \in P$, find the closest point $q_j \in Q$ (Nearest Neighbor).
2. **Estimation:** Find (R, t) that minimize the error for these pairs:

$$(R^*, t^*) = \operatorname{argmin}_{R, t} \sum \|q_j - (Rp_i + t)\|^2$$

3. **Transformation:** Apply $P \leftarrow R^*P + t^*$.
4. **Convergence:** If the change in error $< \epsilon$, stop. Otherwise return to 1.

Note: Step 1 (Nearest Neighbor) is expensive $\rightarrow$ KD-Tree structures.

Interactive demonstration: ICP



● Cible Q (fixe) ● Source P (mobile)

▶ Démarrer

▶ Auto

↺ Reset

The Nearest Neighbor Problem

For each point $p_i \in P$, find the closest point $q_j \in Q$.

Naive solution:

- For each of the n points in P : iterate over the n points in Q and compute the distance.
- Finding one correspondence: $O(n)$
- Finding **all** correspondences: $O(n^2)$ ← too slow!

Scalability problem

A LiDAR produces $\sim 100,000$ points per scan. $O(n^2)$ means 10^{10} operations per ICP iteration – impractical in real time.

K-D Tree: Spatial Data Structure

Efficient solution: organize the points in a **binary spatial partitioning tree**.

Principle:

- Build a binary tree representing the target point cloud Q .
- Traverse the tree instead of a list to find matches.

K-D Tree: Building the Tree

Recursive steps:

1. Choose an axis (alternating $x, y, z, \dots$).
2. Split the point cloud in two based on the **median point** on this axis.
3. Add the median point to the tree (node).
4. Recurse into both halves.

```
Level 0 (X axis): median → root node
```

```
├─ Level 1 (Y axis): left median
```


```
│   └─ Level 2 (X axis): ...
```

```
│       └─ Level 2 (X axis): ...
```

```
└─ Level 1 (Y axis): right median
```

```
    └─ ...
```

Interactive Demonstration: Building a K-D Tree



Prof. 0 Prof. 1 Prof. 2 Prof. 3 Prof. 4

◀ Précédent

Suivant ▶

▶ Auto

↺ Reset

K-D Tree: Construction Complexity

Why $O(n \log n)$?

The tree has $\log_2 n$ levels of depth. At each level, **all n points** are partitioned to find the median:

$$\underbrace{\log n}_{\text{levels}} \times \underbrace{O(n)}_{\text{median}} = O(n \log n)$$

Intuition: recursive splits

n points $\rightarrow$ 1 median	(level 0)
$n/2 + n/2 \rightarrow$ 2 medians	(level 1)
...	
1 point per leaf	(level $\log_2 n$)

At each level, the total work is $O(n)$. The sum over $\log n$ levels gives $O(n \log n)$.

K-D Tree: Nearest Neighbor Search


Search algorithm for a query point q :

1. Starting from the **root**, compute the distance to the current node.
2. Compare q along the node's axis $\rightarrow$ descend into the corresponding branch.
3. At the leaf, keep the best candidate found.
4. **Backtrack up the tree** to check whether a closer point could exist in the other branch (compare the distance to the splitting plane).

Key step: backtracking check

A sphere of radius = current distance may cross the splitting plane $\rightarrow$ the other subtree must be explored if this is the case.

Interactive Demo: Search in the K-D Tree



● Points ● Nœuds visités ● Plus proche voisin ● Point requête

↺ Reset

K-D Tree: Search Complexity

Why $O(\log n)$ on average?

Each comparison on an axis **eliminates** $\sim 1/2$ of the **remaining space**, like a binary search:

$$\underbrace{\log n}_{\text{decisions}} \times O(1) = O(\log n)$$

Operation	Complexity
1 match	$O(\log n)$ avg.
Worst case (high dim.)	$O(n)$
n matches	$O(n \log n)$ avg.

Why $O(n)$ in the worst case?

Going back up the tree, we check whether the sphere of radius = current distance **crosses the separating plane**. In high dimension ($k \gg 3$), this sphere crosses almost all hyperplanes $\rightarrow$ we end up visiting every node.

Gain vs naive solution

$$O(n^2) \xrightarrow{\text{K-D Tree}} O(n \log n)$$

For $n = 100,000$ points:

- Naive: 10^{10} operations
- K-D Tree: $\sim 1.7 \times 10^6$ operations

K-D Tree: Limitations

Advantages:

- Very efficient for representing **points** in low dimension.
- Simple to implement, widely available (Open3D, PCL, scipy).

Limitations:

- For complex geometries (surfaces, volumes), **bounding volume hierarchies** (BVH) are preferred.
- Efficiency **decreases with dimension k** – in high dimension, one must backtrack up the tree very often (*curse of dimensionality*).

Approximation:

Can be built in $O(n \log n)$ using an **approximate** $O(n)$ algorithm to find medians (linear selection algorithm). Requires periodic rebalancing of the tree if the data changes.

In practice

Libraries like **FLANN** or **nanoflann** use approximate variants (ϵ -NN) to further speed up the search.

The Standard ICP Algorithm

Input: P (Source), Q (Target), initial estimate (R_0, t_0) .

1. **Data association:** For each point $p_i \in P$, find the closest point $q_j \in Q$ (Nearest Neighbor).
2. **Estimation:** Find (R, t) that minimize the error for these pairs:

$$(R^*, t^*) = \operatorname{argmin}_{R, t} \sum \|q_j - (Rp_i + t)\|^2$$

3. **Transformation:** Apply $P \leftarrow R^*P + t^*$.
4. **Convergence:** If the change in error $< \epsilon$, stop. Otherwise return to 1.

Note: Step 1 (Nearest Neighbor) is expensive $\rightarrow$ KD-Tree structures.

SVD Resolution: Centering

How to solve step 2 exactly? (Kabsch / Arun method).

1. Computing the centroids:

$$\mu_P = \frac{1}{N} \sum_{i=1}^N p_i, \quad \mu_Q = \frac{1}{N} \sum_{i=1}^N q_i$$

2. Centering the point clouds: We define the points relative to the center of the point cloud:

$$p'_i = p_i - \mu_P, \quad q'_i = q_i - \mu_Q$$

Intuition: This decouples rotation from translation.

SVD Resolution: Rotation

3. Cross-Covariance Matrix:

$$W = \sum_{i=1}^N q'_i (p'_i)^T = Q' P'^T$$

4. Singular Value Decomposition (SVD):

$$W = U D V^T$$

5. Rotation Computation:

$$R = U V^T$$

Intuition: Why does SVD give the rotation?


1. What W contains ($Q'P'^T$)

This matrix captures how the axes of the point cloud P correlate with the axes of the point cloud Q .

2. The role of SVD (UDV^T)

SVD decomposes any linear transformation into:

- V^T : Rotation (aligns P with a canonical basis).
- D : Stretch (scaling).
- U : Rotation (projects onto the basis of Q).



Since we are looking for a **rigid** transformation (no deformation), we ignore the stretch D :

$$R = U \cdot \underbrace{I}_{\text{ignore } D} \cdot V^T = UV^T$$

Solving with SVD: Translation


Once the optimal rotation R is known, the translation is the difference between the centroids, corrected by the rotation.

$$t = \mu_Q - R \mu_P$$

Summary of the estimation step

The SVD approach gives the optimal solution in the least-squares sense for fixed correspondences, without inner iteration.

Interactive Demo: SVD Estimation



Source P (rouge), cible Q (bleu) – correspondances $1 \leftrightarrow 1, 2 \leftrightarrow 2, \dots$

● Source P ● Cible Q 1 / 6

◀ Précédent

Suivant ▶

▶ Auto

↺ Reset

An Important Variant: Point-to-Plane

The "Point-to-Point" metric converges slowly in environments with geometric ambiguities (e.g., corridors).

Point-to-Point:

$$\|q - (Rp + t)\|^2$$

Minimizes the direct Euclidean distance.

Point-to-Plane:

$$\|n_q \cdot (q - (Rp + t))\|^2$$

Minimizes the distance projected onto the normal n_q of the target surface.

- **Advantage:** Allows the cloud to "slide" along planar surfaces, faster convergence and better geometric accuracy.
- **Required:** Surface normals must be computed.

Outlier Management (Aberrant Data)

Naive ICP is very sensitive to outliers (quadratic least squares).

Correspondence rejection strategies:

1. **Max distance threshold:** If $\|p_i - q_j\| > d_{\max}$, the pair is ignored.
2. **Normal compatibility:** If the angle between normals n_p and n_q is too large ($> 45^\circ$), it is not the same surface.
3. **Trimmed ICP:** The errors are sorted and only the best $X\%$ (e.g., 90%) are kept to compute the transformation.

Limitations of ICP

Local convergence problem

- It converges to the nearest local minimum.
- **Consequence:** It requires a good initialization (e.g., odometry, constant velocity).
- **Lab 1:** You will see how matching characteristic points can solve this problem.

Typical applications:

- Scan-to-scan registration (LiDAR odometry).
- Scan-to-map registration (Localization).
- Object tracking.